

Hong Yuan Cao

☎ 857-265-4853 ✉ hongyc@bu.edu 🔗 [linkedin.com/in/hong-yuan-cao/](https://www.linkedin.com/in/hong-yuan-cao/) 🌐 github.com/hongyuanc

EDUCATION

Boston University

Boston, MA

Double Major: Bachelor of Arts in Computer Science and Economics, GPA: 3.74/4.0

Graduation: May 2026

Relevant Courses: AI, ML, Programming Languages, Distributed Systems, Databases, Software Engineering, Analysis of Algorithms, Computer Systems, Probability in Computing, Combinatoric Structures, Discrete Mathematics, Data Structures

TECHNICAL SKILLS

Languages: Python, JavaScript, TypeScript, Java, Go, C, C#, SQL, OCaml, HTML, CSS, Shell, LaTeX, Assembly
Frameworks: Django, React, React Native, Vue.js, Electron, Flask, Node.js, Spring, PyTorch, Pandas, OpenCV, gRPC
Tools: Git, Docker, Unix, Linux, Raspberry Pi, MySQL, PostgreSQL, Redis, AWS, ChromaDB, Terraform, GDB

EXPERIENCES

Software Engineer Intern

Shanghai, China

KeelinMicro (Automotive Startup)

June 2025 – August 2025

- Co-invented a **patent-pending** CV architecture that cuts processing latency by **30-40%** and power usage by **15-20%**
- Engineered an **OpenCV** pipeline with Python to generate 3D point clouds of objects from live video at over **60 FPS**
- Architected a real-time synchronization system for camera and radar inputs, achieving **sub-millisecond accuracy**
- Built custom **Linux kernel driver** in C and **device tree overlay** to enable GMSL camera streaming on Raspberry Pi 5

Software Engineer Intern

Shanghai, China

Tospor Real Estate Consulting Co., Ltd

June 2024 – August 2024

- Built **Django ETL pipeline** on Alibaba Cloud consolidating several databases, reducing data aggregation time by **45%**
- Migrated **200+** consultants from 4 legacy systems to unified platform by implementing secure credential migration APIs
- Optimized database performance via **query caching** and **connection pooling**, reducing query latency by **23%**

Machine Learning Intern

Shanghai, China

Anban Tech (Fuzz Testing Startup)

May 2024 - June 2024

- Built preprocessing pipeline to clean code-testing datasets, improving data quality for ML model training using **Pandas**
- Curated **2.2GB JUnit 5** dataset from open-source repositories using **GitHub API** for fuzzing LLM training
- Developed **Python** validation scripts achieving **95%** accuracy of dataset format validation for training pipeline

Software Engineer Intern

San Francisco, CA

Tonic – Sprout Technologies, Inc. (Entertainment Tech Startup)

June 2023 – October 2023

- Implemented responsive UI components and animations using **React.js**, increasing landing page engagement by **24%**
- Built a social media crawler using **Python (BeautifulSoup4)** that identified **200+** qualified influencers to target
- Filtered **3,000+** media contacts using **Python (Pandas)** and Excel, improving outreach response rates by **25%**

PROJECTS

Athens - AI Learning Platform | *FastAPI, React, TypeScript, OpenAI, Supabase, ChromaDB*

- Built full-stack architecture with **FastAPI**, **React TypeScript**, and hybrid **ChromaDB/Supabase** storage
- Developed end-to-end **RAG** system with document ingestion, vector embeddings, and semantic retrieval for Q&A
- Built **dynamic prompt construction** with context injection and memory, improving response quality by **40%**

Serverless Image Processing Pipeline | *AWS S3, Lambda, CloudFront, DynamoDB, Terraform, React*

- Deployed serverless infrastructure using **Terraform**, **S3**, **Lambda**, **CloudFront**, reducing delivery latency by **38%**
- Implemented **Python**-based **Lambda** functions triggered by **S3** events to automatically resize uploaded images

Multi-Currency Budget Tracking Mobile App | *Django, React Native, PostgreSQL, Android Studio*

- Developed a expense tracking app with **React Native**, featuring currency conversion via third-party API integration
- Built **Django REST APIs** using **composite indexing** and **ORM** optimization, reducing retrieval by **65%**

Movie Recommendation Platform | *Flask, Vue.js, Python, Javascript, PostgreSQL*

- Integrated movie database API with **Flask** to create an interactive dashboard using **Vue.js** components
- Built content-based recommendation algorithm using **CountVectorizer** and **cosine similarity** to analyze metadata

Neural Language Model Implementation | *Python, PyTorch*

- Implemented transformer architecture in **PyTorch** with token and positional embeddings for sequential prediction